

Name: Karen Alfred Habib

ID: 2205236

Social Network Security

Assignment 2

Bot Detection in Facebook Ego Networks

1. Dataset:

- The dataset used is the well-known **Facebook combined graph** from SNAP (Stanford Network Analytics Project), containing :
 - **4,039 nodes** (users) and
 - **88,234 undirected edges** (mutual friendships).
- This is the exact same data family referenced in the assignment link, we used the merged version for a richer and more realistic evaluation.
- We created **realistic synthetic bots** (7% of nodes \approx 282 bots) by selecting originally normal-looking users (degree 20–300) and dramatically reducing their connections to only 3–8 edges each.
- We build a bot-detection classifier using **GraphSAGE**, then we test three scenarios:

- 1. Baseline (no attack)**
- 2. Structural Evasion Attack**
- 3. Graph Poisoning Attack**

Finally, we compare the performance and analyze how each attack affects the model.

2. Dataset & Graph Construction:

We load the Facebook graph using **NetworkX** and build an undirected graph where each node represents a user.

3. Feature Extraction:

For each node, we extract the following graph-based features:

- Node degree
- Clustering coefficient
- PageRank
- Betweenness centrality
- Eigenvector centrality

These features are later used by the ***GraphSAGE*** model.

- Degree & Degree Centrality
- Clustering Coefficient

- Betweenness Centrality
 - Closeness Centrality
 - Eigenvector Centrality
 - Community detection using **Louvain algorithm** (modularity + community size)
-

4. Train/Eval Function (GraphSAGE Model):

- ❖ We split the nodes into:
 - **80% training**
 - **20% testing**
- ❖ We train the GraphSAGE model using:
 - **Adam optimizer**
 - **Negative Log-Likelihood loss**
 - **30 epochs**
- ❖ During training, the model learns how the graph structure relates to bot behavior.
- ❖ After training, we compute:
 - AUC-ROC
 - PR-AUC
 - Precision

- Recall
 - F1-score
-


5. Baseline Scenario (No Attack):

We train and test the model using the original clean graph.

The output from the code:

- ROC ≈ 0.9225
- PR-AUC ≈ 0.8548
- F1-Score ≈ 0.8247

This represents the detector's performance when **no** attacker is modifying the graph.



6. Structural Evasion Attack:

- ❖ In this attack, bot nodes change their connections to “look more normal”.

For example:

- Bots add edges to high-degree real users
- Bots remove links between themselves to avoid being clustered

This makes their graph features closer to humans.

- ❖ After applying the attack, we repeat:

- feature extraction
- model evaluation

❖ **We do NOT retrain the model.**

❖ **Expected behavior/output**

- ROC drops ≈ 0.7873
- PR-AUC drops ≈ 0.2917
- F1 decreases ≈ 0.0351

Because the bots now look more like normal users.

❖ **Code:**

```
#----- Realistic Structural Evasion Attack (Mimicry - FIXED & WORKING) -----
G_evasion = G_baseline.copy()
bots_set_evasion = bots_set_baseline.copy()

print("Applying REALISTIC Structural Evasion Attack – Bots mimic real humans
perfectly")

for bot in bots_set_evasion:
    valid_humans = [n for n in G_evasion.nodes()
                    if n not in bots_set_evasion
                    and 35 <= G_evasion.degree(n) <= 120]
    if not valid_humans:
        continue
    human_template = random.choice(valid_humans)

    #Delete all old links
    G_evasion.remove_edges_from([(bot, neigh) for neigh in
G_evasion.neighbors(bot)])

    # Add the human being's neighbors in an almost identical way
    human_neighbors = list(G_evasion.neighbors(human_template))
    target_neighbors = list(human_neighbors)
    random.shuffle(target_neighbors)
```


```

desired_degree = len(human_neighbors)
num_to_add = random.randint(int(desired_degree * 0.8), int(desired_degree *
1.15))
num_to_add = min(num_to_add, len(target_neighbors))

for target in target_neighbors[:num_to_add]:
    G_evasion.add_edge(bot, target)

print(f"Evasion Attack completed — {len(bots_set_evasion)} bots are now invisible")

```



7. Graph Poisoning Attack:

- ❖ In a poisoning attack, the attacker modifies the training graph before the model is trained.
- ❖ This allows bots to influence the model to learn incorrect patterns.

❖ Usually, poisoning causes:

- Large drop in ROC
- Model failing to detect bots
- Higher false positives

❖ Expected behavior/output “Training: After Poisoning (Retrained)”:

- ROC drops ≈ 0.9371
- PR-AUC drops ≈ 0.8460
- F1 decreases ≈ 0.7020

❖ Code:

```
# ----- Graph Poisoning Attack -----
G_poison = G_baseline.copy()
bots_set_poison = bots_set_baseline.copy()


print("Injecting 200 medium-degree poison bots...")
for i in range(200):
    new_bot = max(G_poison.nodes()) + 1 + i
    G_poison.add_node(new_bot)

    # poison bots with human-like degree → confuse the model
    deg = random.randint(20, 60)
    targets = random.sample(list(G_poison.nodes()), deg)

    for t in targets:
        G_poison.add_edge(new_bot, t)

    bots_set_poison.add(new_bot)

print("Poisoning completed")
```



8. After-Poisoning Test on Clean Data:

We also test the poisoned model on the **clean graph** again.

This shows how much the poisoning corrupted the detector.

❖ Expected result

- Huge decrease in performance
- Bot detection becomes unreliable

❖ **Expected behavior/output “Evaluation Only (NO retraining):
Poisoned Model → Clean Data”:**

- ROC drops ≈ 0.8911
 - PR-AUC drops ≈ 0.6784
 - F1 decreases ≈ 0.5195
-
-

Main observations:

- Structural evasion makes bots “blend in”, so performance drops moderately.
 - Poisoning attacks corrupt the learning process, so the detector fails badly.
 - After poisoning, the model becomes unreliable even on clean data.
-

9. Conclusion:

This assignment shows that bot detection based only on graph structure is vulnerable to adversarial attacks.

- **Structural evasion** makes bots hide inside the network.
- **Poisoning attacks** are even more dangerous because they corrupt the training process.
- To improve robustness, future models should include:

- adversarial training
 - anomaly detection
 - temporal behavior analysis
-